



INFO20003 Database Systems

Xiuge Chen

Tutorial 6
2020.04.26



- 1. Storage and Indexing review (key concepts with examples) - 25 min**
- 2. More SQL - 10min**
- 3. Exercises - 25min**
- 4. Lab - 50 min**

Storage and Indexing

Why?

- Database management systems store information on disks (normally hard disks)
- involves many READ and WRITE operations when data is accessed: **high cost**
- **READ**: transfer of data from the disk to main memory (RAM)
- **WRITE**: transfer data from RAM to the disk



Storage and Indexing

alternative terms used with respect to disk storage

Conceptual modelling	Entity	Attribute	Instance of an entity
Logical modelling	Relation	Attribute	Tuple
Physical modelling/SQL	Table	Column/Field	Row
Disk storage	File	Field	Record




Storage and Indexing

Files, pages and records

- **record**: an individual **row** of a table and has a unique *rid* (disk address of the page containing the record).
ex. *rid* (3, 7) refers to the seventh record from third page
- **page**: an allocation of space on disk or in memory containing a collection of **records**. (every page is the **same size**)
- **file**: a collection of **pages** containing records (In simple database scenarios: single table)

Storage and Indexing

Files, pages and records





Storage and Indexing

File organisations


- defines how file records are mapped onto pages (stored on disk).
- Heap file organisation:
- Sorted file organisation
- Index file organisation



Storage and Indexing

Heap file organisation:

- No ordering, sequencing or indexing
- Suitable when retrieving all records
- Slow search
- Quick insert



Storage and Indexing

Sorted file organisation:

- sequential order based on the *search key* (not PK or FK)
- Quick search for search key (especially on range)
- Slow insert

Page 1

record 8

record 9

record 10

Page 2

record 12

record 16

record 21

Page 3

record 33

record 45

Page 4

Storage and Indexing

Index file organisation:

- Any **subset** of the fields of a table is indexed based on queries that are frequently run against the database
- Quick search on the subset attributes
- Different index could be built
- Different types index could be chosen
- Insert depend on types

Storage and Indexing

What is index?

- made up of **data entries** which refer back to the data in the relation. (k, rid) k : search key, rid : record ID.
- **speeds up** selection on the search key fields
- **search key**: subset of the attributes of a relation on which the index is built (not be relation's key!!!)
- stored in an **index file**, in contrast to the **data file** which contains the actual records themselves

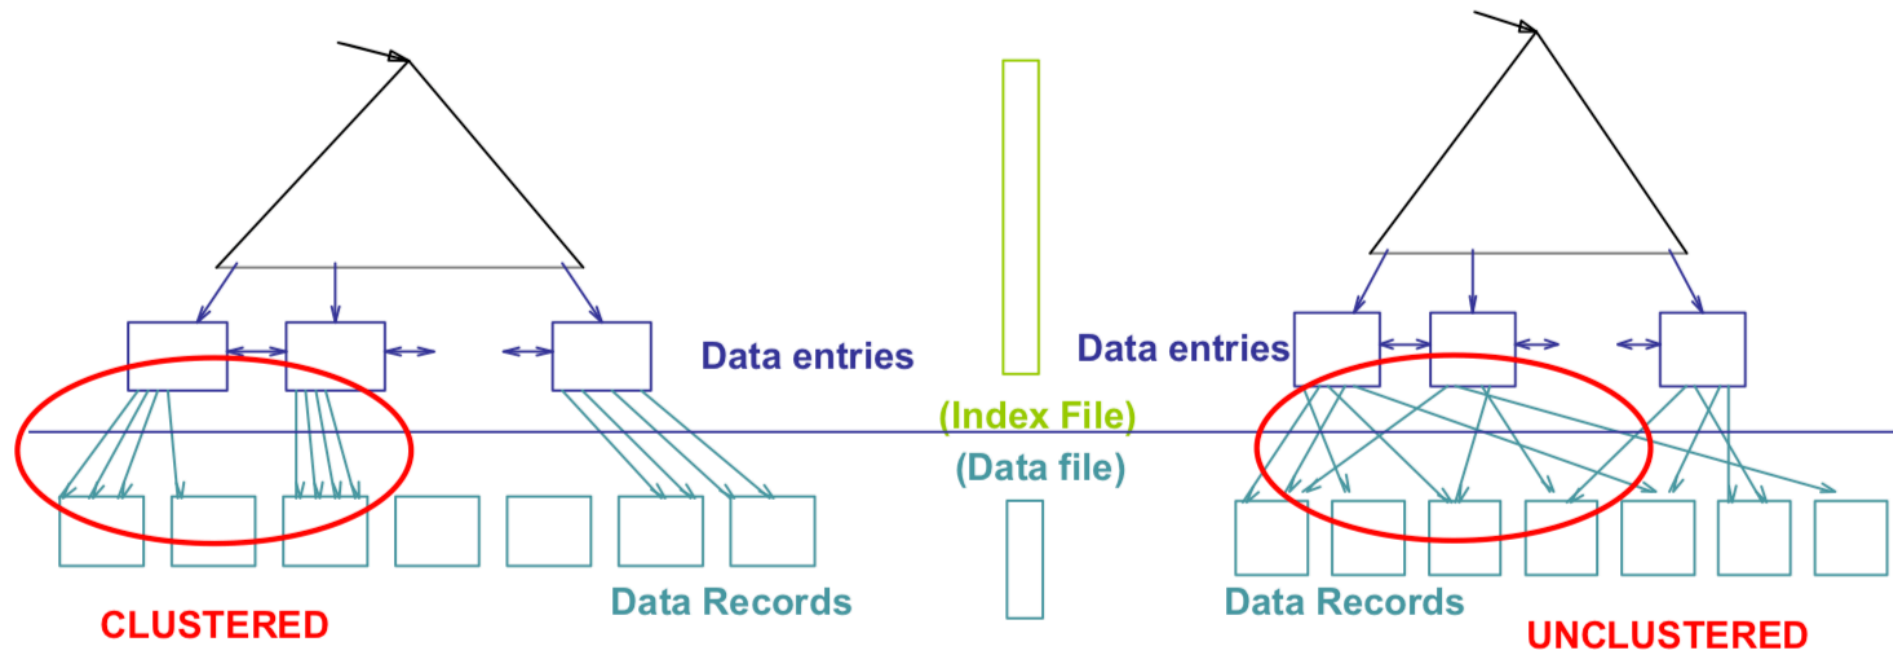
Storage and Indexing

Type of index 1

- **clustered**: data records in the data file have the **same order** as data entries of the index
- **unclustered**: data records in the data file are **not sorted by** search key/data entries
- **primary**: on the primary key of the relation
- **secondary**: on any other set of attributes


Storage and Indexing

Type of index 1




Storage and Indexing

Type of index 1



Storage and Indexing

Type of index 1



Storage and Indexing

How to choose index

- Which relations are accessed frequently?
- Which attributes are retrieved?
- Which attributes are involved in selection, join and other conditions?
- If a query involves updating the relation, what attributes are affected?

Storage and Indexing

Index

- Coming: how to analyze a given query plan and see if a better query plan exists with an additional index
- **In general:** make SELECT queries faster but slow down the updates
- Indexes also require **additional disk space**.
- carefully analyzed before constructing an index!!!

Storage and Indexing

Type of index 2

- **Hash-based indexing**: hash function $h(r)$ is applied, where r is the field value.
- **Output**: point to a bucket which refers to the primary page and other overflow pages if there is any. These buckets contain a representation (k, rid) for data entries.
- best suited to support *equality* selections
- How to build: not in this subject

Storage and Indexing

Hash index

- Suppose you are given 5 buckets and $h(k) = k \% 5$ where % is the modulus (remainder) operator. Insert 200, 22, 119, 8, and 33 into a hash table.

Bucket	Key
0	200
1	
2	22
3	8, 33
4	119



Storage and Indexing

Hash index

- <https://www.cs.usfca.edu/~galles/visualization/OpenHash.html>

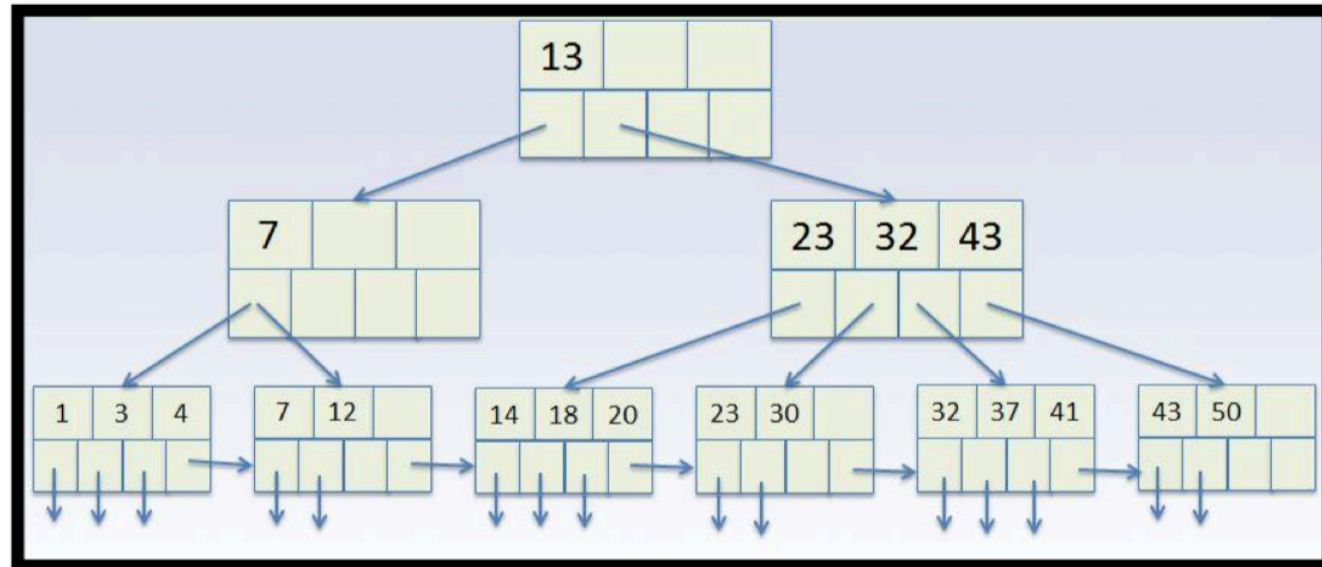
Storage and Indexing

Type of index 3

- **B-tree index**: sorting data on search key and maintaining a hierarchical search data structure (B+ tree) that will direct the search to the respective page of the data entry
- **Insertion** in such a structure is **costly** as the tree is updated with every insertion or deletion.
- Good for: **equal or range selections**
- How to build: not in this subject

Storage and Indexing

B-tree index



- Start at the root.
- Internal nodes, search the keys to find the range K belongs in and follow that pointer.
- For leaf nodes (nodes with no child nodes), search the keys to find K and follow the pointer to find the data record.



Storage and Indexing

B-tree index

- <https://www.cs.usfca.edu/~galles/visualization/BPlusTree.html>



Any questions?

Structured Query Language(SQL)

Operator	Description
=	Equal to
<	Less than
>	Greater than
<=	Less than or equal to
>=	Greater than or equal to
<> OR !=	Not equal to (depends on DBMS as to which is used)

Structured Query Language(SQL)

- **Logic:** AND, NOT, OR
- [UNION] ALL: do not eliminate duplicates
- IN / NOT IN
- AS
- SUM / COUNT / MAX / MIN
- YEAR ...

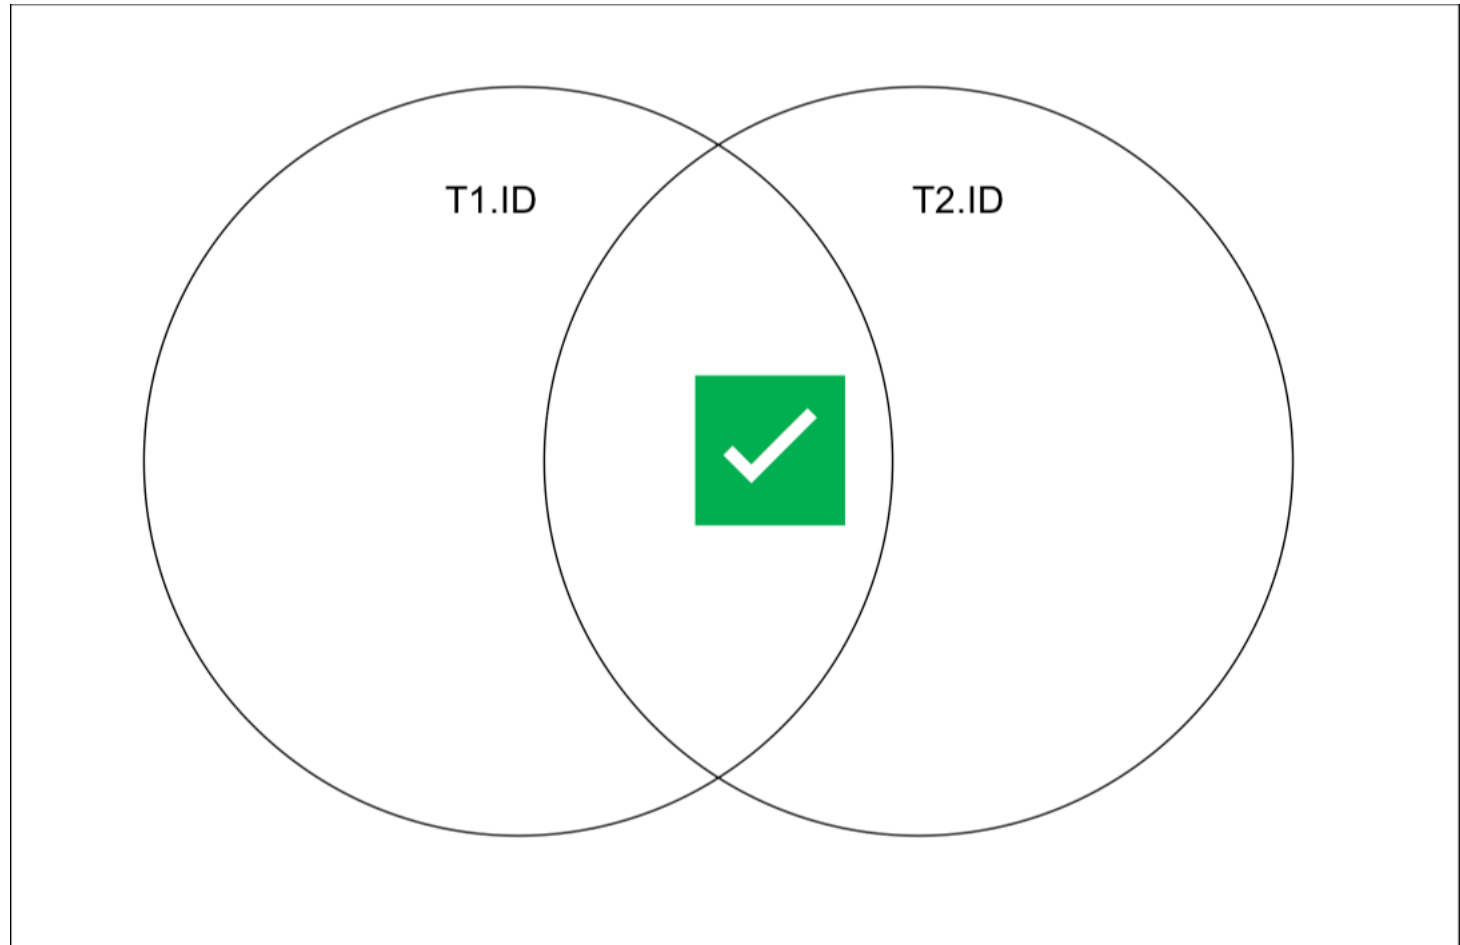
Structured Query Language(SQL)

Inner Join

INNER

JOIN

ON




Structured Query Language(SQL)

Left Outer
Join

LEFT
OUTER
JOIN

ON

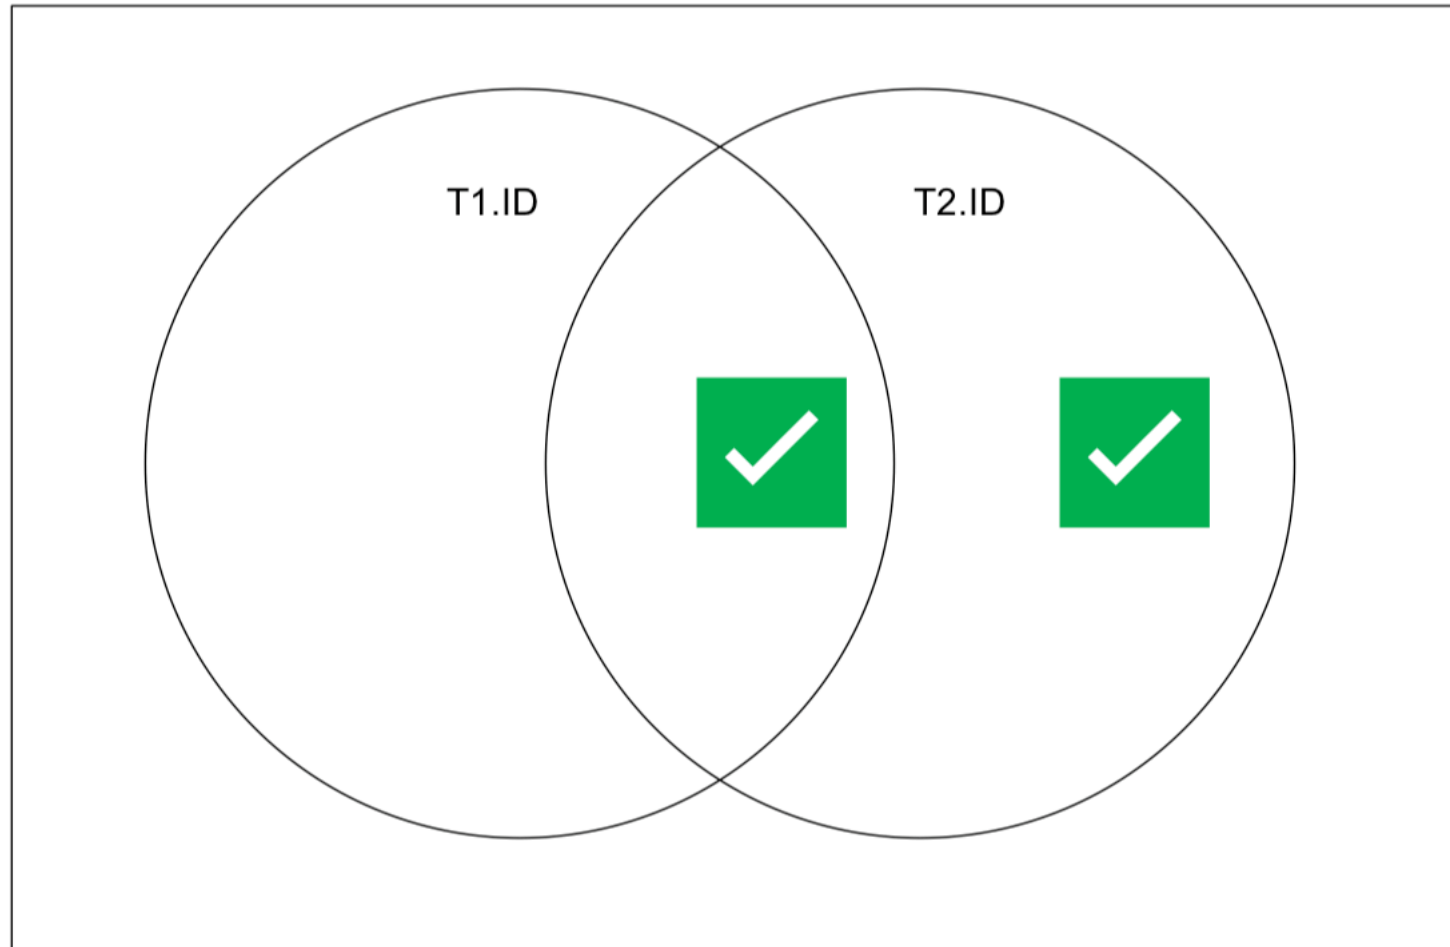


Structured Query Language(SQL)

Right Outer
Join

RIGHT
OUTER
JOIN

ON




Structured Query Language(SQL)

Full Outer
Join

FULL
OUTER
JOIN

ON



Structured Query Language(SQL)

- **SELECT** [ALL | DISTINCT] *select_expr* [, *select_expr* ...]
 - List the columns (and expressions) that are returned from the query
- **[FROM *table_references*]**
 - Indicate the table(s) or view(s) from where the data is obtained
- **[WHERE *where_condition*]**
 - Indicate the conditions on whether a particular row will be in the result
- **[GROUP BY {*col_name* | *expr* } [ASC | DESC], ...]**
 - Indicate categorisation of results
- **[HAVING *where_condition*]**
 - Indicate the conditions under which a particular category (group) is included in the result
- **[ORDER BY {*col_name* | *expr* | *position*} [ASC | DESC], ...]**
 - Sort the result based on the criteria
- **[LIMIT {[*offset*,] *row_count* | *row_count* OFFSET *offset*}]**
 - Limit which rows are returned by their return order (ie 5 rows, 5 rows from row 2)



Any questions?

1. Choosing an index

You are asked to create an index on a suitable attribute. What are the important aspects you will analyse to make this decision? To get you started, the following might help you by providing scaffolding to the discussion:

a. Primary vs. secondary index

Primary: records are retrieved based on the value of primary key.

Secondary: fields that are frequently queried.

Generally, a table should always have a primary index (in fact, MySQL creates one automatically).

b. Clustered vs. unclustered index

Clustered: consists of a frequently-executed condition to check for a **range**, however expensive to maintain

Unclustered: fields that are frequently queried.

Equality conditions: same if the search key does not have duplicate values.

More than one combination of columns is used in range queries, choose the **most frequently** used combination and make those fields search keys of the clustered index



c. Hash vs. tree indexes

Hash: equality queries, faster than B-tree

Tree: range queries, creating a B-tree index



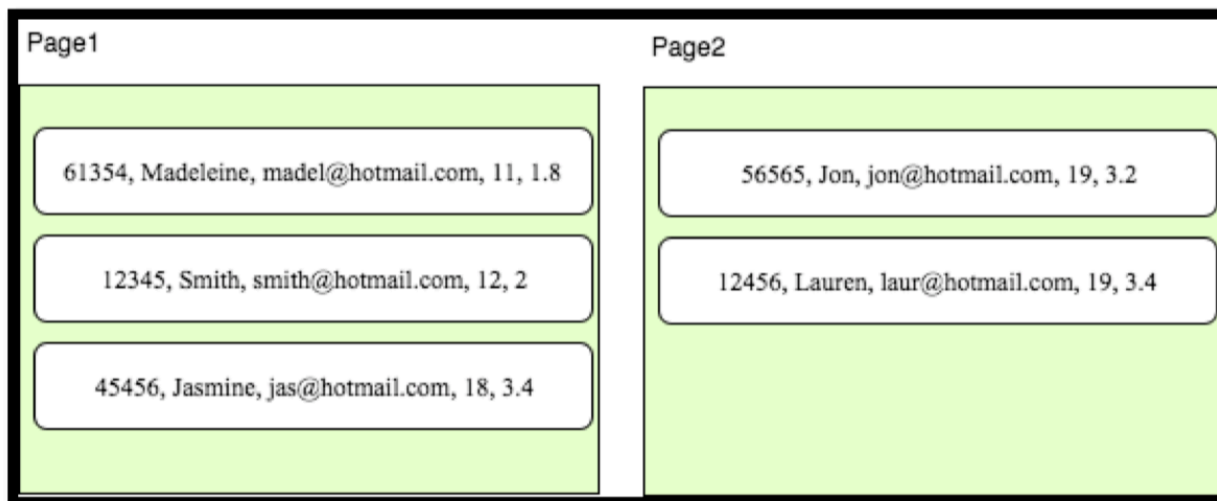
1. Data entries of an index:

SID	Name	Email	Age	GPA
61354	Madeleine	madel@hotmail.com	11	1.8
12345	Smith	smith@hotmail.com	12	2.0
45456	Jasmine	jas@hotmail.com	18	3.4
56565	Jon	jon@hotmail.com	19	3.2
12456	Lauren	laur@hotmail.com	19	3.4

1. tuples sorted by age

2. order of tuple is the same when stored on disk

3. each page can contain only 3 records



1. Data entries of an index:

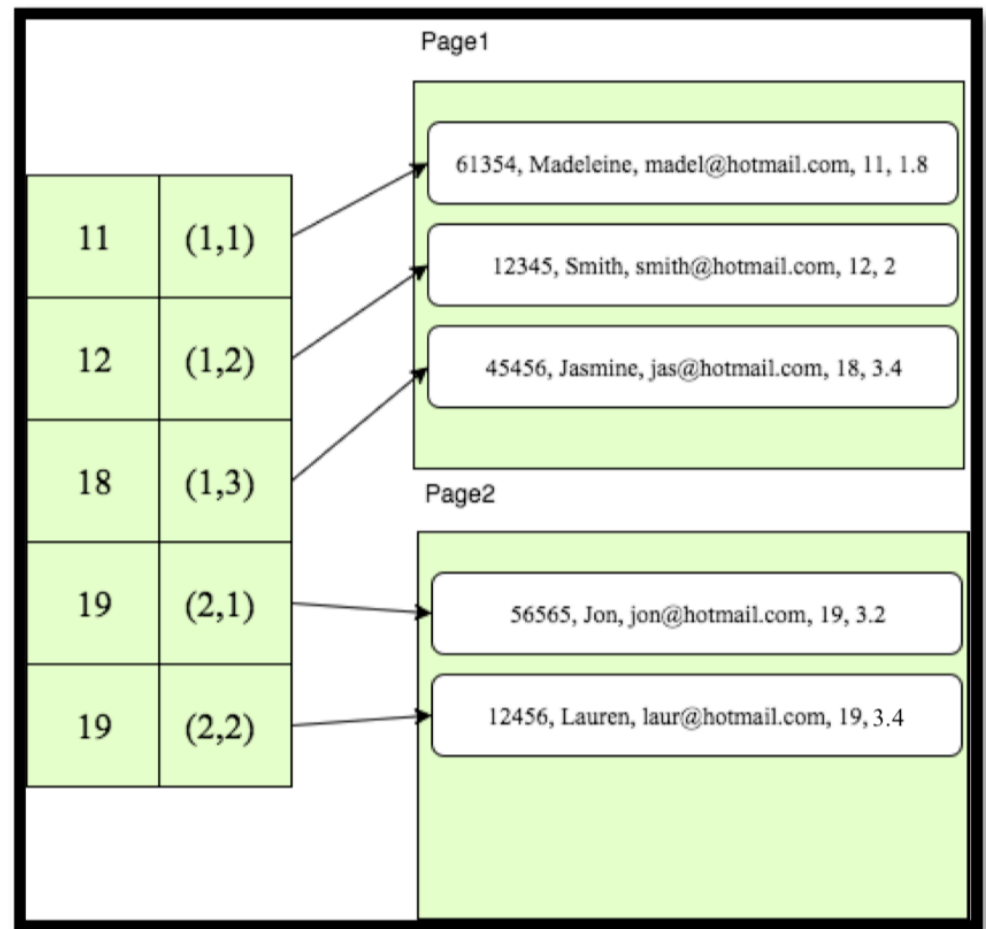
Show what the *data entries* of the index will look like for:

a. An index on Age

search key and *rid* in the format (a, b)

a is the page number and
 b is the record number.

clustered index



1. Data entries of an index:

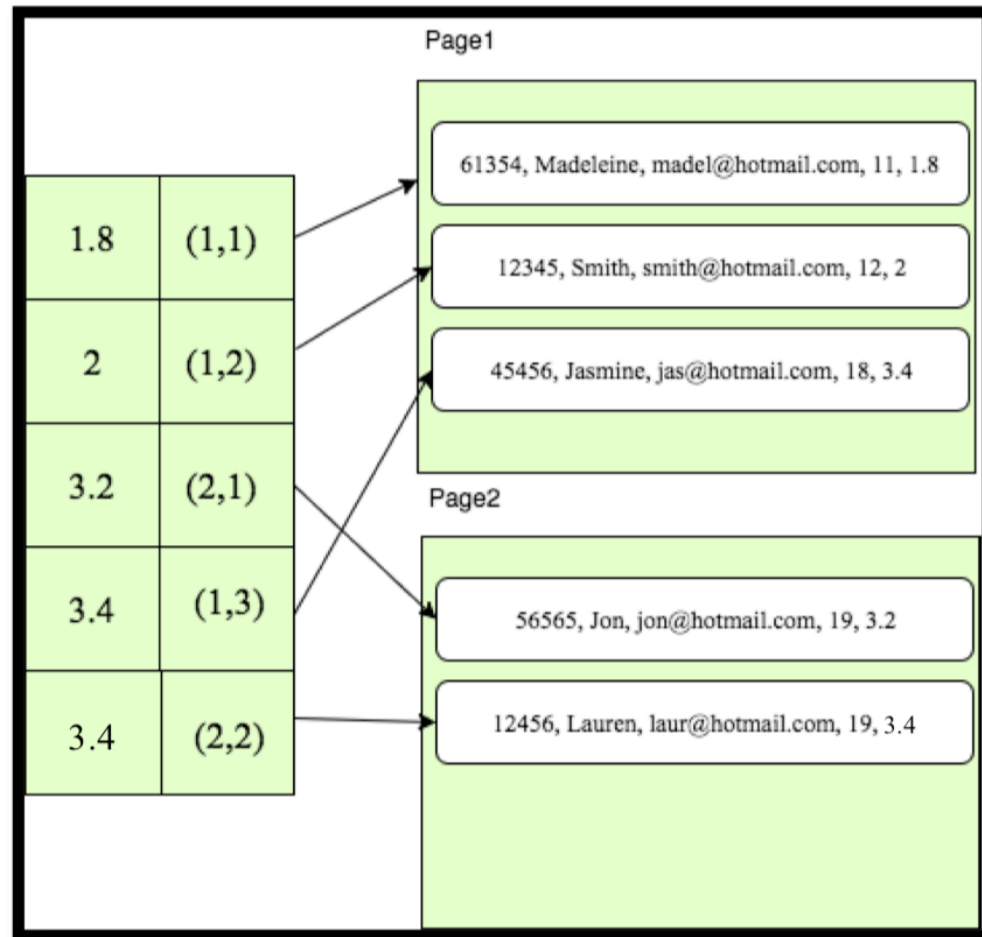
Show what the *data entries* of the index will look like for:

b. An index on GPA

search key and *rid* in the format (a, b)

a is the page number and
 b is the record number.

unclustered index



1. Consider the following relations:

Employee (EmployeeID, EmployeeName, Salary, Age, DepartmentID)^{FK}

Department (DepartmentID, DepartmentBudget, DepartmentFloor, ManagerID)^{FK}

In the database, the salary of employees ranges from AUD10,000 to AUD100,000, age varies from 20-80 years and each department has 5 employees on average. In addition, there are 10 floors, and the budgets of the departments vary from AUD10,000 to AUD 1million.

Given the following two queries frequently used by the business, which index would you prefer to speed up the query? Why?

1. Consider the following relations:

Employee (EmployeeID, EmployeeName, Salary, Age, DepartmentID) ^{FK}

Department (DepartmentID, DepartmentBudget, DepartmentFloor, ManagerID) ^{FK}

a. **SELECT** DepartmentID
FROM Department
WHERE DepartmentFloor = 10
AND DepartmentBudget < 15000;

- A) Clustered hash index on DepartmentFloor
- B) Unclustered hash Index on DepartmentFloor
- C) Clustered B+ tree index on (DepartmentFloor, DepartmentBudget)
- D) Unclustered hash index on DepartmentBudget
- E) No need for an index



1. Consider the following relations:

Employee (EmployeeID, EmployeeName, Salary, Age, DepartmentID) ^{FK}

Department (DepartmentID, DepartmentBudget, DepartmentFloor, ManagerID) ^{FK}

a. **SELECT** DepartmentID
FROM Department
WHERE DepartmentFloor = 10
AND DepartmentBudget < 15000;

- A) Clustered hash index on DepartmentFloor
- B) Unclustered hash Index on DepartmentFloor
- C) Clustered B+ tree index on (DepartmentFloor, DepartmentBudget)
- D) Unclustered hash index on DepartmentBudget
- E) No need for an index

Range query!

1. Consider the following relations:

Employee (EmployeeID, EmployeeName, Salary, Age, DepartmentID) ^{FK}

Department (DepartmentID, DepartmentBudget, DepartmentFloor, ManagerID) ^{FK}

b. **SELECT** EmployeeName, Age, Salary
FROM Employee;

- A) Clustered hash index on (EmployeeName, Age, Salary)
- B) Unclustered hash index on (EmployeeName, Age, Salary)
- C) Clustered B+ tree index on (EmployeeName, Age, Salary)
- D) Unclustered hash index on (EmployeeID, DepartmentID)
- E) No need for an index



1. Consider the following relations:

Employee (EmployeeID, EmployeeName, Salary, Age, DepartmentID) ^{FK}

Department (DepartmentID, DepartmentBudget, DepartmentFloor, ManagerID) ^{FK}

b. **SELECT** EmployeeName, Age, Salary
FROM Employee;

- A) Clustered hash index on (EmployeeName, Age, Salary)
- B) Unclustered hash index on (EmployeeName, Age, Salary)
- C) Clustered B+ tree index on (EmployeeName, Age, Salary)
- D) Unclustered hash index on (EmployeeID, DepartmentID)
- E) No need for an index

get requested attributes with an **index-only scan** (and we can avoid accessing the table completely)



Any questions?



Please refer to Lab 6 on LMS

Let me know if you encounter with
any problem

More practice on SQL Skills